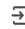


```
!pip install -q datasets
```

```
!pip install --upgrade sentence-transformers==4.1.0
```

 [Show hidden output](#)

```
import transformers, sentence_transformers
print(f"transformers: {transformers.__version__}")
print(f"sentence-transformers: {sentence_transformers.__version__}")
```

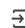
 transformers: 4.51.3
sentence-transformers: 4.1.0

Should be: transformers: 4.51.3 sentence-transformers: 4.1.0

```
from huggingface_hub import notebook_login
notebook_login()
```

```
from datasets import load_dataset
```

```
dataset = load_dataset("aegean-ai/ai-lectures-spring-24")
print(dataset) # Inspect the dataset structure
```

 DatasetDict({
 train: Dataset({
 features: ['mp4', 'info.json', 'en.vtt', 'json', '__key__', '__url__'],
 num_rows: 8
 })
})

Double-click (or enter) to edit

✓ Encode embeddings

```
!pip install fuzzywuzzy
!pip install bertopic
```

 [Show hidden output](#)

```
import cv2
import numpy as np
from PIL import Image
import tempfile
import os
import re
from collections import OrderedDict
import spacy
from spacy import displacy
from bertopic import BERTopic
```

```
from datasets import load_dataset
import json
import os
# Create a directory to save the JSON files if it doesn't exist
os.makedirs("json_files", exist_ok=True)

# Access the train split (since that's what your dataset shows)
train_dataset = dataset['train']

# Iterate through each example and save the JSON file
for i, example in enumerate(train_dataset):
    # Get the JSON data
    json_data = example['json']

    # Define the output filename (you can customize this)
    output_filename = f"json_files/lecture_{i}.json"

    # Save the JSON data to a file
    with open(output_filename, 'w') as f:
        json.dump(json_data, f, indent=2) # indent for pretty printing
```

```
print(f"Saved {len(train_dataset)} JSON files to the 'json_files' directory.")
```

 Saved 8 JSON files to the 'json_files' directory.

```
sample = dataset['train']
# subtitles = sample["en.vtt"] # Subtitles are in VTT format
metadata = sample["json"]
```

```
def advanced_clean(text):
    """Specialized cleaning for lecture transcripts"""
    # Remove filler words and artifacts
    text = re.sub(r'\b(uh|um|ah|kind of|sort of|you know)\b', '', text, flags=re.IGNORECASE)
    # Remove XML-like tags and timestamps
    text = re.sub(r'<[>]+>', '', text)
    # Remove repeated phrases
    text = re.sub(r'(\b\w+\b)(?:\s+\1)+', r'\1', text)
```

```

# Remove very short sentences
if len(text.split()) < 5:
    return ""
return text.strip()

def remove_repeated_phrases(text):
    # This regex finds repeated phrases (1 to ~15 words) and removes duplicates
    pattern = r'\b((\w+\s+){1,15})\1'
    return re.sub(pattern, r'\1', text, flags=re.IGNORECASE)

def remove_substring_sentences(sentences):
    sentences = list(set(sentences)) # Remove exact duplicates first
    filtered = []

    for s in sentences:
        if not any((s != other and s in other) for other in sentences):
            filtered.append(s)

    return filtered

def convert_to_seconds(timestamp):
    """Convert various timestamp formats to seconds (float)"""
    if isinstance(timestamp, (int, float)):
        return float(timestamp)
    elif isinstance(timestamp, str):
        if ':' in timestamp: # Format like "00:01:30.500"
            parts = timestamp.split(':')
            if len(parts) == 3: # HH:MM:SS.sss
                return float(parts[0]) * 3600 + float(parts[1]) * 60 + float(parts[2])
            elif len(parts) == 2: # MM:SS.sss
                return float(parts[0]) * 60 + float(parts[1])
            return float(timestamp) # Try direct conversion
        return 0.0 # Default fallback

nlp = spacy.load("en_core_web_sm")

def process_vtt_metadata(metadata):
    chunks = [] # Will contain lists of sentences for each video
    time_chunks = [] # Will contain timestamp information for each sentence
    all_caption_texts = []
    all_timestamps = [] # To store timestamp info for each caption

    # Loop through each video in the dataset
    for video in metadata:
        captions = video['captions']
        for caption in captions:
            all_caption_texts.append(caption['text'])
            # Store start and end times for each caption
            all_timestamps.append((caption['start'], caption['end']))

        single_line_text = " ".join(all_caption_texts)
        single_line_text = advanced_clean(single_line_text)
        single_line_text = remove_repeated_phrases(single_line_text)

        # Process with spaCy
        doc = nlp(single_line_text)

        # Get sentences and approximate their timestamps
        sentences = [sent.text for sent in doc.sents]
        sentence_timestamps = approximate_sentence_timestamps(sentences, all_timestamps, all_caption_texts)

        chunks.append(sentences)
        time_chunks.append(sentence_timestamps)

        # Reset for next video
        all_caption_texts = []
        all_timestamps = []

    return chunks, time_chunks

from fuzzywuzzy import fuzz # or: from thefuzz import fuzz

def approximate_sentence_timestamps(sentences, caption_timestamps, caption_texts):
    """
    Approximate timestamps for sentences using fuzzy matching with original captions.
    All timestamps are converted to and handled in seconds.
    """
    sentence_times = []

    # Convert all caption timestamps to seconds upfront
    caption_timestamps_seconds = [
        (convert_to_seconds(start), convert_to_seconds(end))
        for start, end in caption_timestamps
    ]

    for sentence in sentences:
        # Find best matching caption for start of sentence
        start_match_idx, start_score = find_best_fuzzy_match(
            sentence, caption_texts, is_start=True
        )
        start_time = caption_timestamps_seconds[start_match_idx][0] if start_match_idx != -1 else None

        # Find best matching caption for end of sentence
        end_match_idx, end_score = find_best_fuzzy_match(
            sentence, caption_texts, is_start=False
        )
        end_time = caption_timestamps_seconds[end_match_idx][1] if end_match_idx != -1 else None

```

```

# Fallback logic if no good matches found
if start_time is None or end_time is None:
    if not sentence_times: # First sentence
        start_time = caption_timestamps_seconds[0][0]
        end_time = caption_timestamps_seconds[0][1]
    else:
        prev_end = sentence_times[-1][1] # Already in seconds
        avg_duration = max(3.0, len(sentence.split()) * 0.5) # More dynamic duration estimate
        start_time = prev_end
        end_time = prev_end + avg_duration
# Ensure start comes before end
elif start_time >= end_time:
    end_time = start_time + max(1.0, len(sentence.split()) * 0.3)

sentence_times.append((start_time, end_time))

return sentence_times

def find_best_fuzzy_match(sentence, captions, is_start=True, threshold=75):
    """
    Find the best fuzzy match between a sentence and captions.
    If is_start=True, looks at the beginning of the sentence.
    If is_start=False, looks at the end of the sentence.
    """
    best_idx = -1
    best_score = 0

    # Take first/last 10 words for matching
    words = sentence.split()
    sample_size = min(10, len(words))
    query = " ".join(words[:sample_size] if is_start else words[-sample_size:])

    for i, caption in enumerate(captions):
        # Use partial ratio since we're matching parts of texts
        score = fuzz.partial_ratio(query.lower(), caption.lower())

        if score > best_score:
            best_score = score
            best_idx = i

    # Only return if above threshold
    return (best_idx, best_score) if best_score >= threshold else (-1, best_score)

```

▼ process Embeddings

```

a, b = process_vtt_metadata(metadata)

import pickle
# Save to file
def save_data(data, filename):
    with open(filename, 'wb') as f: # 'wb' = write binary
        pickle.dump(data, f)

# Load from file
def load_data(filename):
    with open(filename, 'rb') as f: # 'rb' = read binary
        return pickle.load(f)

save_data((a, b), 'caption_data.pkl')

# Load saved data
loaded_chunks, loaded_time_chunks = load_data('caption_data.pkl')

from sentence_transformers import SentenceTransformer
clip_model = 'clip-ViT-B-32'
# sentence_model = SentenceTransformer("all-MiniLM-L6-v2")
sentence_model = SentenceTransformer(clip_model)

🔗 Using a slow image processor as `use_fast` is unset and a slow processor was saved with this model. `use_fast=True` will be the default behavior

total_models = []
all_embeddings = [] # To store embeddings for each video

for video in loaded_chunks:
    # Get embeddings first
    video_embeddings = sentence_model.encode(video, show_progress_bar=True)
    all_embeddings.append(video_embeddings) # Store embeddings

    # Then fit BERTopic
    topic_model = BERTopic(embedding_model=sentence_model)
    topics, probs = topic_model.fit_transform(video)
    total_models.append(topic_model)
    print(topic_model.get_topic_info())

```



Batches: 100% 1/1 [00:00<00:00, 1.17it/s]

Topic	Count	Name \
0	-1	12 -1_to_and_the_we

Representation \

0 [to, and, the, we, of, that, is, you, are, in]

Representative_Docs

0 [in this video I would like to start the discu...

Batches: 100% 2/2 [00:00<00:00, 7.68it/s]

Topic	Count	Name \
0	0	38 0_this_so_that_of
1	1	17 1_the_we_this_and

Representation \

0 [this, so, that, of, and, in, the, okay, be, we]

1 [the, we, this, and, of, that, is, in, to, will]

Representative_Docs

0 [so that is we the zero and we'll be getting i...

1 [and maybe the I mean if right and maybe th...

Batches: 100% 5/5 [00:00<00:00, 13.07it/s]

Topic	Count	Name \
0	0	77 0_the_of_we_to
1	1	69 1_okay_so_the_one

Representation \

0 [the, of, we, to, that, and, is, this, in, be]

1 [okay, so, the, one, and, right, this, is, be,...

Representative_Docs

0 [so if I may and as far as the output feature...

1 [okay, okay, okay]

Batches: 100% 2/2 [00:00<00:00, 22.21it/s]

Topic	Count	Name \
0	0	30 0_so_and_we_to
1	1	18 1_the_we_of_and

Representation \

0 [so, and, we, to, this, have, definitely, our,...

1 [the, we, of, and, to, in, have, is, that, are]

Representative_Docs

0 [and so that's basically our head set, and I...

1 [so here is the point where the head so here i...

Batches: 100% 1/1 [00:00<00:00, 14.83it/s]

Topic	Count	Name \
0	0	11 0_the_we_and_of
1	1	11 1_let_say_64_the

Representation \

0 [the, we, and, of, are, that, this, to, see, so]

1 [let, say, 64, the, those, so, have, of, compo...

Representative_Docs

0 [in an earlier video we saw the structure of t...

1 [so this is what we have seen we can this is ...

Batches: 100% 4/4 [00:00<00:00, 13.95it/s]

Topic	Count	Name \
0	0	100 0_the_so_of_this
1	1	24 1_the_we_of_to

Representation \

0 [the, so, of, this, okay, to, is, y0, y1, we]

1 [the, we, of, to, and, that, in, have, so, is]

Representative_Docs

0 [so the second term is simply F52 of y1 y 0 + ...

1 [I think it's worthwhile providing some guid...

Batches: 100% 6/6 [00:00<00:00, 14.23it/s]

Topic	Count	Name \
0	0	75 0_it_so_to_and
1	1	58 1_the_and_we_to
2	2	28 2_the_is_so_to

Representation \

0 [it, so, to, and, going, here, we, okay, that,...

1 [the, and, we, to, is, that, of, this, in, have]

2 [the, is, so, to, what, of, this, just, and, we]

Representative_Docs

0 [so we always going to have present here, so i...

1 [definitely this is a very easy to evaluate t...

2 [so what we will of every predictor so what we...

Batches: 100% 2/2 [00:00<00:00, 16.92it/s]

Topic	Count	Name \
0	-1	8 -1_is_of_the_linear
1	0	28 0_so_this_is_the
2	1	26 1_the_to_of_we

Representation \

0 [is, of, the, linear, to, so, and, discriminat...

1 [so, this, is, the, we, of, and, it, generaliz...

2 [the, to, of, we, is, and, this, probability, ...

Representative_Docs

0 [so this is going to be effectively to be ef...

1 [and this is the probability of X and this is ...

2 [x this is the form of U logistic regression s...

Final Embeddings are located in all_embeddings

The original sentences are located in "loaded_chunks" The time stamps for each sentence are in "loaded_time_chunks"

▼ Insert into Qdrant

```
!pip install -q moviepy qdrant-client
from sentence_transformers import SentenceTransformer
from moviepy.editor import VideoFileClip
from google.colab import files
import os

from qdrant_client.models import VectorParams, Distance

from qdrant_client import QdrantClient, models

# --- Corrected Video Segment Upload Script ---

from qdrant_client import QdrantClient
from qdrant_client.models import Distance, VectorParams, PointStruct
from sentence_transformers import SentenceTransformer
from uuid import uuid4
import numpy as np

# --- Start Qdrant client (in memory)
client = QdrantClient(":memory:")

collection_name = "video_captions"

# --- Load embedding model
embedding_model = SentenceTransformer("all-MiniLM-L6-v2") # 384 dimensions
print("Embedding model loaded.")

# --- Generate real embeddings
real_embeddings = []
for sentences in loaded_chunks: # loaded_chunks = list of list of sentences
    sentence_embeddings = embedding_model.encode(sentences)
    real_embeddings.append(sentence_embeddings)

# --- Detect embedding dimension
embedding_dim = real_embeddings[0].shape[1] # Correctly detect 384
print(f"Detected embedding dimension: {embedding_dim}")

# --- Recreate collection with correct size
try:
    client.get_collection(collection_name)
    print(f"Collection '{collection_name}' already exists.")
except:
    client.create_collection(
        collection_name=collection_name,
        vectors_config=VectorParams(
            size=embedding_dim,
            distance=Distance.COSINE,
        )
    )
    print(f"Created collection '{collection_name}' with {embedding_dim} dimensions.")

# --- Insert video segments into Qdrant
for video_idx, (sentences, embeddings, time_data) in enumerate(zip(loaded_chunks, real_embeddings, loaded_time_chunks)):
    points = []

    # Handle time_data correctly: it's a list of (start, end) pairs
    starts = [start for start, end in time_data]
    ends = [end for start, end in time_data]

    # Verify data alignment
    assert len(sentences) == len(embeddings) == len(starts) == len(ends)

    for sentence_idx, (sentence, embedding, start, end) in enumerate(zip(
        sentences,
        embeddings,
        starts,
        ends
    )):
        points.append(
            PointStruct(
                id=str(uuid4()),
                vector=embedding.tolist(), # Convert numpy array to list
                payload={
                    "video_id": video_idx,
                    "text": sentence,
                    "position": sentence_idx,
                    "start_sec": float(start),
                    "end_sec": float(end),
                    "duration_sec": float(end) - float(start),
                }
            )
        )
    client.upsert(collection_name, points)
```

```
# Upload batch to Qdrant
operation_info = client.upsert(
    collection_name=collection_name,
    points=points,
    wait=True
)

print(f"✅ Uploaded {len(points)} segments for Video {video_idx} "
      f"(Time range: {starts[0]:.1f}s to {ends[-1]:.1f}s)")
```

```
🔗 Embedding model loaded.
Detected embedding dimension: 384
Created collection 'video_captions' with 384 dimensions.
✅ Uploaded 12 segments for Video 0 (Time range: 3.0s to 286.7s)
✅ Uploaded 55 segments for Video 1 (Time range: 2.0s to 906.5s)
✅ Uploaded 146 segments for Video 2 (Time range: 3.3s to 2392.0s)
✅ Uploaded 48 segments for Video 3 (Time range: 666.5s to 666.6s)
✅ Uploaded 22 segments for Video 4 (Time range: 2.3s to 503.1s)
✅ Uploaded 124 segments for Video 5 (Time range: 281.2s to 281.2s)
✅ Uploaded 161 segments for Video 6 (Time range: 1.9s to 1.9s)
✅ Uploaded 62 segments for Video 7 (Time range: 5.6s to 1603.8s)
```

✓ Retrieval Processing

```
# Install Python wrapper
!pip install imageio[ffmpeg]
!pip install ffmpeg-python
```

```
# Install system ffmpeg
!apt-get update
!apt-get install ffmpeg
```

```
🔗 Requirement already satisfied: imageio[ffmpeg] in /usr/local/lib/python3.11/dist-packages (2.37.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from imageio[ffmpeg]) (2.0.2)
Requirement already satisfied: pillow>=8.3.2 in /usr/local/lib/python3.11/dist-packages (from imageio[ffmpeg]) (11.1.0)
Requirement already satisfied: imageio-ffmpeg in /usr/local/lib/python3.11/dist-packages (from imageio[ffmpeg]) (0.6.0)
Requirement already satisfied: psutil in /usr/local/lib/python3.11/dist-packages (from imageio[ffmpeg]) (5.9.5)
Requirement already satisfied: ffmpeg-python in /usr/local/lib/python3.11/dist-packages (0.2.0)
Requirement already satisfied: future in /usr/local/lib/python3.11/dist-packages (from ffmpeg-python) (1.0.0)
Hit:1 https://cloud.r-project.org/bin/linux/ubuntu jammy-cran40/ InRelease
Get:2 http://security.ubuntu.com/ubuntu jammy-security InRelease [129 kB]
Hit:3 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2204/x86\_64 InRelease
Hit:4 https://r2u.stat.illinois.edu/ubuntu jammy InRelease
Hit:5 http://archive.ubuntu.com/ubuntu jammy InRelease
Get:6 http://archive.ubuntu.com/ubuntu jammy-updates InRelease [128 kB]
Hit:7 https://ppa.launchpadcontent.net/deadsnakes/ppa/ubuntu jammy InRelease
Hit:8 https://ppa.launchpadcontent.net/graphics-drivers/ppa/ubuntu jammy InRelease
Hit:9 https://ppa.launchpadcontent.net/ubuntugis/ppa/ubuntu jammy InRelease
Get:10 http://archive.ubuntu.com/ubuntu jammy-backports InRelease [127 kB]
Fetched 384 kB in 2s (210 kB/s)
Reading package lists... Done
W: Skipping acquire of configured file 'main/source/Sources' as repository 'https://r2u.stat.illinois.edu/ubuntu jammy InRelease' does not seem
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
ffmpeg is already the newest version (7:4.4.2-0ubuntu0.22.04.1).
0 upgraded, 0 newly installed, 0 to remove and 35 not upgraded.
```

✓ Gradio APP using FFMPEG

```
# utils/video_utils.py
import ffmpeg
import numpy as np
from io import BytesIO

def extract_segment(video_bytes, start_sec, end_sec):
    """Extract video segment using FFmpeg without saving to disk."""
    try:
        out, _ = (
            ffmpeg.input('pipe:0')
            .trim(start=start_sec, end=end_sec)
            .setpts('PTS-STARTPTS') # Reset timestamps
            .output('pipe:', format='mp4', vcodec='libx264')
            .run(input=video_bytes, capture_stdout=True, capture_stderr=True)
        )
        return out
    except ffmpeg.Error as e:
        print(f"FFmpeg error: {e.stderr.decode()}")
        return None
```

```
!pip install gradio
```

```
🔗 Show hidden output
```

```
# --- HEALTH CHECK for Qdrant and Dataset ---
```

```
# Test 1: Check if Qdrant collection exists
try:
    collections = client.get_collections()
    print(f"✅ Qdrant collections found: {collections}")
except Exception as e:
    print(f"❌ Error connecting to Qdrant: {e}")
```

```
# Test 2: Try searching manually (fallback to .search())
```

```

try:
    dummy_vec = embedding_model.encode("test query").tolist()
    hits = client.search(
        collection_name=collection_name,
        query_vector=dummy_vec,
        limit=1
    )
    print(f"✅ Retrieved {len(hits)} hits from Qdrant.")
except Exception as e:
    print(f"❌ Error during Qdrant search: {e}")

# Test 3: Check if dataset has mp4 bytes
try:
    sample = dataset[0]
    if "mp4" in sample and isinstance(sample["mp4"], bytes):
        print("✅ Dataset has valid mp4 bytes.")
    else:
        print("❌ Dataset missing mp4 bytes.")
except Exception as e:
    print(f"❌ Error accessing dataset: {e}")

⚠️ WARNING:py.warnings:<ipython-input-35-62d387b4f00b>:13: DeprecationWarning: `search` method is deprecated and will be removed in the future. Use
hits = client.search(

✅ Qdrant collections found: collections=[CollectionDescription(name='video_captions')]
✅ Retrieved 1 hits from Qdrant.
❌ Error accessing dataset: "Invalid key: 0. Please first select a split. For example: `my_dataset_dictionary['train'][0]`. Available splits: ['

# --- Retrieve top subtitles given a user query ---

def retrieve_top_subtitles(question, top_k=3):
    try:
        # 1. Embed the user's question
        query_embedding = embedding_model.encode(question).tolist()

        # 2. Search in Qdrant
        hits = client.search(
            collection_name=collection_name,
            query_vector=query_embedding,
            limit=top_k
        )

        # 3. Collect the subtitles
        subtitles = []
        for hit in hits:
            if hasattr(hit, "payload") and "subtitles" in hit.payload:
                subtitle_text = hit.payload["subtitles"]
                subtitles.append(subtitle_text)

        if subtitles:
            print(f"✅ Retrieved {len(subtitles)} subtitles:")
            for i, subtitle in enumerate(subtitles, 1):
                print(f"Top {i}: {subtitle}")
        else:
            print("⚠️ No subtitles found in hits.")

        return subtitles

    except Exception as e:
        print(f"❌ Error during retrieval: {e}")
        return []

sample = dataset[0]

# Step 1: See available keys
print(f"Sample keys: {sample.keys()}")

# Step 2: See what type 'mp4' is
print(f"Type of sample['mp4']: {type(sample['mp4'])}")

# Step 3: Preview only first 100 characters/bytes
if isinstance(sample["mp4"], bytes):
    print(f"First 100 bytes: {sample['mp4'][:100]}")
elif isinstance(sample["mp4"], str):
    print(f"First 100 characters: {sample['mp4'][:100]}")
else:
    print("Unexpected format inside 'mp4' field.")

```

 [Show hidden output](#)

Next steps: [Explain error](#)

```

# --- Corrected Gradio App with Subtitles Summary ---

import gradio as gr
import tempfile
import os
from datasets import load_dataset # Ensure you have datasets installed
# NOTE: Make sure you already initialized client, collection_name, embedding_model separately!

# --- Load your HuggingFace dataset ---
dataset = load_dataset("aegean-ai/ai-lectures-spring-24", split="train")

# --- Helper function ---
def extract_segment(video_bytes, start_sec, end_sec):
    temp_dir = tempfile.mkdtemp()

```

```

temp_path = os.path.join(temp_dir, f"clip_{start_sec:.2f}_{end_sec:.2f}.mp4")

with open(temp_path, "wb") as f:
    f.write(video_bytes)

return temp_path

# --- Main retrieval + answering function ---
def answer_with_video(question):
    # 1. Encode the question
    question_embedding = embedding_model.encode(question).tolist()

    # 2. Qdrant search (use .search because of old client version)
    hits = client.search(
        collection_name=collection_name,
        query_vector=question_embedding,
        limit=3
    )

    # --- Safety: No hits found ---
    if not hits:
        return "No relevant video segments found.", "", "", ""

    # --- Extract video clips and subtitles ---
    video_clip_paths = []
    subtitle_texts = []

    for i, hit in enumerate(hits):
        video_idx = hit.payload.get("video_id", 0)
        start = hit.payload.get("start_sec", 0)
        end = hit.payload.get("end_sec", 5)

        # Safety: check if video_idx is valid
        if video_idx >= len(dataset):
            return f"Error: Invalid video_idx {video_idx} returned.", "", "", ""

        sample = dataset[video_idx]

        # Safety: check mp4 field
        if "mp4" not in sample or not isinstance(sample["mp4"], bytes):
            return "Error: mp4 bytes not found in dataset.", "", "", ""

        video_bytes = sample["mp4"]

        if video_bytes is None:
            return "Error: No video bytes available.", "", "", ""

        # Save video bytes temporarily
        temp_dir = tempfile.gettempdir()
        clip_path = os.path.join(temp_dir, f"clip_{i}.mp4")

        with open(clip_path, "wb") as f:
            f.write(video_bytes)

        video_clip_paths.append(clip_path)

        # Collect subtitle if available
        subtitle = hit.payload.get("subtitles", "")
        if subtitle:
            subtitle_texts.append(subtitle)

    # --- Generate better answer text ---
    if subtitle_texts:
        answer_text = "Summary based on retrieved clips:\n\n" + "\n".join(
            [f"- {text}" for text in subtitle_texts]
        )
    else:
        answer_text = f"Found {len(video_clip_paths)} relevant segments."

    # --- Ensure 3 outputs ---
    while len(video_clip_paths) < 3:
        video_clip_paths.append("")

    return answer_text, video_clip_paths[0], video_clip_paths[1], video_clip_paths[2]

# --- Gradio App ---
with gr.Blocks() as demo:
    gr.Markdown("## 🎥 Chat with Your Course Videos (RAG System)")

    with gr.Row():
        question = gr.Textbox(label="Ask a question about the course videos")
        submit_btn = gr.Button("Search")

    answer = gr.Textbox(label="Generated Answer")
    video1 = gr.Video(label="Relevant Clip 1")
    video2 = gr.Video(label="Relevant Clip 2")
    video3 = gr.Video(label="Relevant Clip 3")

    submit_btn.click(
        fn=answer_with_video,
        inputs=question,
        outputs=[answer, video1, video2, video3]
    )

# Launch the app
demo.launch(share=True)

```


Colab notebook detected. To show errors in colab notebook, set debug=True in launch()
* Running on public URL: <https://f7de51f44494364414.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working directory to

Chat with Your Course Videos (RAG System)

Ask a question about the course videos

Search

Using only the videos, explain how ResNets work.

Generated Answer

Found 3 relevant segments.

